

From Concern Detection to Promotion Control: A Protocol for Fail-Closed Governance of Autonomous Research Artifacts

Anonymous

Abstract

Autonomous research systems increasingly produce complete-looking manuscripts and run bundles, yet review components can identify serious scientific defects without preventing a paper-ready decision. We formulate this gap as *promotion governance*: deciding whether an artifact bundle should be promoted, reviewed, downgraded, or blocked, while linking each blocking concern to inspectable evidence and a responsible repair stage. We introduce a manifest-driven evaluation protocol with clean controls, counterfactual fault variants, base-bundle-disjoint splits, blind manuscript-only requests, and metrics for false promotion, concern-acceptance conflict, clean-case retention, repair localization, traceability, and cost. A synthetic development suite with four base bundles and 40 cases validates the evaluator and exposes three missing integrity checks in an artifact audit. After those checks were added, the audit passed all development cases; an advisory-only ablation still accepted every defective case despite detecting its concerns. These results are instrument validation, not empirical evidence for general effectiveness: the suite is synthetic, unreviewed, and explicitly ineligible for paper-level claims. We specify the independently adjudicated, real-provider experiment needed to test the central hypothesis.

1 Introduction

Research agents can search literature, implement experiments, analyze results, and draft papers [1]. Capability benchmarks therefore evaluate increasingly complete replication and research workflows [2]. Completion, however, is not equivalent to scientific readiness. BadScientist shows that plausible but unsound papers can receive favorable automatic reviews [3], while counterfactual evaluation finds that automatic reviewers often fail to alter their judgments when faulty reasoning is injected [4].

We study the policy boundary between detecting a concern and allowing an artifact bundle to advance. A useful research governor must fail closed: a blocking concern cannot coexist with a paper-ready decision unless the concern is resolved or explicitly downgraded. It must also avoid the trivial solution of blocking everything. This motivates three questions: whether artifact-bound gates reduce false promotion, whether binding concerns to transitions prevents concern-acceptance conflict, and what clean-case, repair-localization, latency, and cost trade-offs follow.

The contribution at the present evidence level is a protocol and tested instrumentation, not a confirmed performance claim. We provide (i) a portable artifact-bundle task with explicit promotion and repair labels, (ii) leakage and adjudication gates, (iii) baselines that isolate detection from action binding, and (iv) a development failure analysis that routes audit misses back to the responsible workflow stages.

2 Related Work

Research-agent capability. Agent Laboratory organizes literature review, experimentation, and reporting within a multi-agent research workflow [1]. PaperBench evaluates replication of published machine-learning research from paper understanding through executed artifacts [2]. These systems motivate artifact-level evaluation, but their primary object is research capability rather than the consistency of a promotion decision with detected blockers.

Review and claim verification. BadScientist operationalizes concern–acceptance conflict for unsound generated papers [3]. Dycke and Gurevych inject reasoning faults to test whether automatic reviews respond appropriately [4]. CLAIM-BENCH evaluates claim–evidence reasoning in AI papers [5]. Our task retains concern detection but adds an explicit action contract, artifact evidence references, and a repair owner.

Artifact assurance and recovery. ARA reconstructs scientific workflows for reproducibility assessment [6]; REPRO-Bench evaluates agents using papers and replication packages [7]. MADS-CPS proposes machine-checkable run-level admissibility for autonomous laboratories [8]. SAGE routes experiment failures to structured intervention levels [9], and ClaimGarden controls claim states under evidence updates [10]. The closest distinction is therefore not artifact inspection or failure attribution alone, but whether detected scientific blockers deterministically constrain paper-readiness promotion without over-blocking clean bundles.

3 Promotion Governance Protocol

3.1 Task and labels

Each case is a portable research-run artifact tree paired with a private case manifest. The evaluated system returns one decision from `promote`, `needs_review`, `downgrade`, or `block`; zero or more concerns with severity and evidence paths; and the workflow stages responsible for repair. Gold labels remain outside the visible artifact tree.

Every suite declares an evidence class, adjudication status, and paper-claim eligibility. A paper-eligible suite must be independently double adjudicated. All variants derived from one base bundle remain in the same split, and source content hashes detect renamed duplicates across splits.

3.2 Fault families and controls

The protocol requires clean positive, null, and negative controls together with variants covering missing comparison evidence, repeated-run provenance, hidden failed execution, declared-versus-executed budget mismatch, result–figure conflict, claim–evidence conflict, citation support mismatch, stale persisted state, and unsupported claim strength. Each mutation records before and after hashes. Human adjudicators must verify that a mutation introduces only its declared fault.

3.3 Comparison conditions

We define four deterministic conditions and one provider condition. *Always promote* trusts the incoming ready state. A *presence checklist* verifies required files without semantic consistency checks. An *advisory artifact audit* reports the same concerns, evidence, and repair owners as the full audit but does not bind them to the decision. The *fail-closed artifact audit* applies the binding policy. Finally, a blind *manuscript-only reviewer* receives only paper text under opaque request identifiers; case labels, mutations, paths, and gold data remain private.

3.4 Metrics and inference

Primary metrics are false paper-ready promotion and concern–acceptance conflict. We also report four-way decision macro-F1, clean-case promotion, blocker precision/recall/F1, exact repair-owner match, trace coverage, latency, and provider cost. Every system trial must cover every case. Pairwise effects are computed per case, with 5,000-replicate bootstrap intervals clustered by base bundle and exact paired sign tests over base-bundle effects. Synthetic development suites remain exploratory regardless of numeric performance.

4 Implementation and Development Validation

The protocol is implemented as manifest-driven builders, runners, scorers, and blind request import/export commands. Paths are confined to suite roots; symbolic links, duplicate cases, hash mismatches, split

Condition	Acc.	False prom.	Conflict	Clean	Blocker F1
Always promote	.10	1.00	–	1.00	–
Presence checklist	.10	1.00	–	1.00	–
Advisory audit	.10	1.00	1.00	1.00	1.00
Fail-closed audit	1.00	.00	.00	1.00	1.00

Table 1: Synthetic development results over four base bundles and 40 cases. The labels were generated for instrument debugging and were not independently adjudicated; the table is not confirmatory evidence.

leakage, and incomplete trial coverage are validation failures. Failure analysis emits review diagnostics and node-strengthening recommendations with explicit recheck conditions.

The development corpus contains four synthetic base bundles, one clean control and nine counterfactual variants per bundle, for 40 cases. It is marked as synthetic development evidence with unreviewed adjudication and no paper-claim eligibility. Table 1 therefore tests instrumentation only.

Before strengthening, the fail-closed audit reached .70 decision accuracy and .333 false-promotion rate, missing repeated-run provenance, executed-budget, and stale-state faults. The failure analyzer routed the first two to experiment execution and the third to review. Added integrity checks were then evaluated on the same frozen suite. In the final development run, the fail-closed audit differed from the advisory ablation by .90 decision accuracy and -1.00 false promotion. The cluster bootstrap intervals are degenerate because each synthetic fault repeats identically across only four base bundles; the paired sign-test value is .125. This is direct evidence that the development sample is too small for inferential claims, despite perfect case-level scores.

5 Confirmatory Study and Readiness Gate

The frozen confirmatory study requires at least 20 independently sourced base bundles and 200 held-out cases, with 30 base bundles preferred. Two reviewers must adjudicate held-out labels and mutation isolation. The manuscript-only condition requires three preserved real-provider trials per case. At least one post-repair rerun per fault family must measure recovery and regression. Raw requests, responses, decisions, failures, manifests, hashes, latency, and cost must be retained.

The central hypothesis is that fail-closed governance lowers false promotion by at least 20 absolute percentage points relative to the presence checklist while retaining at least 90 percent clean-case promotion. The study is stopped or downgraded if the full policy blocks more than 10 percent of clean bundles, if the checklist is within five points on both primary safety and clean-case metrics, if mutations leak labels, or if independent adjudication is absent.

6 Limitations and Ethical Considerations

Counterfactual faults may be easier than naturally occurring research defects, and deterministic checks can overfit a known artifact schema. A promotion gate does not establish scientific truth; it enforces consistency between available evidence, stated concerns, and allowed transitions. Provider-based reviewers introduce model, prompt, and temporal dependence. Public artifact bundles must exclude credentials, private paths, personal data, and proprietary research content. False promotion can amplify unsound claims, while false blocking can suppress valid null or negative results; both errors must remain visible.

7 Conclusion

Promotion governance treats paper readiness as an evidence-bound transition rather than a by-product of workflow completion. The implementation and development suite show that this distinction is measurable and that detected failures can drive targeted audit strengthening. They do not yet show that the policy generalizes. Independent artifact curation, adjudication, real-provider review, and a frozen confirmatory run are required before advancing beyond a research memo.

References

- [1] Samuel Schmidgall, Yusheng Su, Ze Wang, Ximeng Sun, Jialian Wu, Xiaodong Yu, Jiang Liu, Michael Moor, Zicheng Liu, and Emad Barsoum. Agent laboratory: Using LLM agents as research assistants. In Christos Christodoulopoulos, Tanmoy Chakraborty, Carolyn Rose, and Violet Peng, editors, *Findings of the Association for Computational Linguistics: EMNLP 2025*, pages 5977–6043, Suzhou, China, November 2025. Association for Computational Linguistics.
- [2] Giulio Starace, Oliver Jaffe, Dane Sherburn, James Aung, Jun Shern Chan, Leon Maksin, Rachel Dias, Evan Mays, Benjamin Kinsella, Wyatt Thompson, Johannes Heidecke, Amelia Glaese, and Tejal Patwardhan. Paperbench: Evaluating ai’s ability to replicate ai research, 2025. arXiv version v3, accessed 2026-07-16.
- [3] Fengqing Jiang, Yichen Feng, Yuetai Li, Luyao Niu, Basel Alomair, and Radha Poovendran. BadScientist: Can a research agent write convincing but unsound papers that fool LLM reviewers? In Maria Liakata, Viviane P. Moreira, Jiajun Zhang, and David Jurgens, editors, *Proceedings of the 64th Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 24712–24727, San Diego, California, United States, July 2026. Association for Computational Linguistics.
- [4] Nils Dycke and Iryna Gurevych. Automatic reviewers fail to detect faulty reasoning in research papers: A new counterfactual evaluation framework. *Transactions of the Association for Computational Linguistics*, 14:465–488, 2026.
- [5] Shashidhar Reddy Javaji, Yupeng Cao, Haohang Li, Yangyang Yu, Nikhil Muralidhar, and Zining Zhu. Can AI validate science? benchmarking LLMs on claim $\rightarrow$ Evidence reasoning in AI papers. In Kentaro Inui, Sakriani Sakti, Haofen Wang, Derek F. Wong, Pushpak Bhattacharyya, Biplab Banerjee, Asif Ekbal, Tanmoy Chakraborty, and Dharendra Pratap Singh, editors, *Proceedings of the 14th International Joint Conference on Natural Language Processing and the 4th Conference of the Asia-Pacific Chapter of the Association for Computational Linguistics*, pages 2355–2379, Mumbai, India, December 2025. The Asian Federation of Natural Language Processing and The Association for Computational Linguistics.
- [6] Kevin Riehl, Andres L. Marin, Nikofors Zacharof, Fan Wu, Patrick Langer, Robert Jakob, Anastasios Kouvelas, Georgios Fontaras, and Michail A. Makridis. Ara: Agentic reproducibility assessment for scalable support of scientific peer-review, 2026. arXiv version v2, accessed 2026-07-16.
- [7] Chuxuan Hu, Liyun Zhang, Yeji Lim, Aum Wadhvani, Austin Peters, and Daniel Kang. Repro-bench: Can agentic ai systems assess the reproducibility of social science research?, 2025. arXiv version v1, accessed 2026-07-16.
- [8] Mateo Petel. MADS-CPS: A machine-checkable admissibility contract for AI scientists in autonomous laboratories. In *ICML 2026 Workshop on AI for Science*, 2026.
- [9] Jie Ma, Binfei Chu, Jie Gao, Jinlu Zhang, Yiwei Ma, Yi Tan, Jiayi Ji, Xiaoshuai Sun, and Rongrong Ji. One reflection is not enough: Self-correcting autonomous research via multi-hypothesis failure attribution, 2026. arXiv version v1, accessed 2026-07-16.
- [10] Hafumi Nishi. ClaimGarden: Update-aware claim-state control for AI scientist workflows. In *ICML 2026 Workshop on AI for Science*, 2026.